

FRAME


Brand URL to editable short-form video

Start with the working demo

[Open Frame](#)

1. Paste a brand URL and follow the live generation stages.
2. Play the finished video and inspect its source-linked research.
3. Compare available drafts, edit the copy or swap a reaction, then rebuild and download.

No login required. Generation takes a few minutes; one job runs at a time. Existing examples are also available to browse.



What happens under the hood

In plain English	Implementation
Understand the brand	Website extraction and structured brand profile
Find useful audience context	DDGS brand/category searches, retries and excerpt checks
Write and improve the story	Six drafts, AI critique and revision of the winner
Check the alternate direction	Cited reference plus a separate alignment review
Choose a reaction and assemble	Seven-clip catalogue, mood/situation ranking, Pillow + FFmpeg
Review and download	Editable text/caption; silent 1080 x 1920 H.264 MP4

Validation and practical limits

36 automated tests passed. Live URL runs, video decoding and public page/video delivery were checked. Research uses indexed excerpts, not a complete social feed or measured trend growth. Music and posting are not implemented. Reaction clips are personal-demo references.

Temporary access: this public tunnel depends on the host computer staying awake and online. The URL may change if the tunnel is restarted.

Rationale

Format, implementation and tradeoffs

I prioritised the long-overlay format highlighted in the brief: a relatable scene, a specific problem and a dry payoff. The reaction sits above a separate reading panel so the face and copy stay clear. The trial scope favoured one complete, editable workflow over several unfinished formats.

Python handles website extraction and orchestration. Claude, through the Anthropic SDK, builds the brand profile, writes six candidates and critiques the strongest. DDGS supplies indexed social references; evidence and alignment checks support a second direction when suitable sources exist. Pillow and FFmpeg produce the final video. Claude Code and Codex assisted development.

The main constraints were the trial timeframe, a seven-clip reaction library and uneven access to useful website and social content. Thin websites can weaken the profile. Search excerpts do not establish trend growth, and AI editorial scores do not prove audience performance. Exports are silent; the demo clips are not a cleared commercial asset library.

What we use now vs. how we would scale

Area	Current prototype	Next step for scale
Hosting	Local Python server and temporary public tunnel.	Persistent cloud service, stable domain and access controls.
Rendering	One queued job at a time; files stored locally.	Durable job queue, render workers, object storage and CDN delivery.
Research	DDGS indexed excerpts with retries and a 24-hour cache.	Reliable platform or licensed data feeds; track changes over time to measure momentum.
Media	Seven curated reactions; silent exports in one format.	Expand licensed, labelled footage; add music and slideshow rendering.
Quality	AI critique, source checks and manual review.	Stronger fact checks, creator feedback and performance data to calibrate selection.
Operations	Single-host demo using API calls for generation.	Usage limits, cost monitoring, failure alerts and automatic job recovery.