

FRAME

Brand URL in, on-brand short-form video out

Try it

Demo: <https://ships-intent-convicted-anaheim.trycloudflare.com/>

Code: <https://github.com/qaisak/hardlaunch-task>

1. Paste any brand URL. No login. About two minutes, progress shown on the page.
2. Play the video. Read the brand profile and the sources it found.
3. Swap the reaction, edit the text, rebuild (no new AI call), download the MP4.

The demo runs from my laptop through a temporary link, so it depends on that machine staying online.



The pipeline today, step by step

Nine steps. Three are Claude calls that need judgement; the rest is Python. Every step writes its result to disk before the next one runs, so a failure late in the run never loses the work before it.

Step	What happens	How it is checked
1. Ingest	Download the homepage plus up to two about pages, strip to text.	Under 1,500 characters means a script-rendered or blocked site: fall back to Claude web fetch. Under 400 after that: stop with a clear error.
2. Brand profile	Claude turns the text into fixed fields: niche, audience, core pain, the tools and habits people use instead, scenes, tone, facts allowed, things not to say.	Inferred fields are labelled as inferred. Numbers are only allowed if the site states them.
3. Research	Two short search phrases, then DDGS search restricted to Reddit, TikTok and Instagram. Cached 24 hours.	Claude keeps a result only if it describes a first-hand frustration and quotes an exact substring of the snippet as evidence. Search failure falls back to evergreen only.
4. Format	Claude recommends overlay, slideshow or hook+demo with a reason.	Recorded, not acted on. Only overlay is rendered in this build.
5. Six drafts	Six overlays in the house style, each a different scene and hook, brand named once.	Style rules are a plain text file: observational opening, the no-X list, escalation, deadpan verdict, 30 to 70 words, market-matched slang.
6. Critique	A strict editor scores each draft on on-brand, voice, hook, not-an-ad, comment bait, honesty and length, and names one fix.	Honesty scores 1 if a draft claims anything the profile does not support, which removes it from contention.
7. Revise	The fix is applied to the winner. If references exist, a second trend version is written.	Revise is optional: if it fails the draft ships unrevised. Trend text must stay 30 to 70 words and never claim a trend is live.
8. Clip match	Read the mood of the text, rank the clip library by mood and situation, show the reasons.	No match means a hard error, never a random clip. A person can override the choice.
9. Render	Text drawn into a reading panel under the clip, looped to reading time, ffmpeg H.264, 1080 x 1920.	Poster is extracted from the actual MP4. Provenance (clip, mood, duration, source) saved next to it.

Every Claude call passes through one helper that retries once with a JSON-repair prompt if the output does not parse. 36 automated tests cover clip matching, missing footage, reading time, long-text layout, trend and discovery validation.

Rationale

Format, tools, tradeoffs

Format. Long overlay, because it is your dominant format, it is text-first so it generalises to any URL in four hours, and your reference piece gives a voice that can be written down as rules. One complete, editable format beat three unfinished ones.

Tools. Python; Claude Opus 5 through the Anthropic SDK (including its web fetch tool as the scrape fallback); requests and BeautifulSoup; DDGS for search; Pillow and ffmpeg for the render; a stdlib HTTP server for the studio. Built with Claude Code and Codex.

Tradeoffs. One rendered format. Seven placeholder reaction clips, not a cleared library. Editor scores are model opinion, not audience data. Search excerpts are inspiration, not proof of a trend. Silent exports. Thin websites give thin profiles.

The pipeline at scale

What the same idea looks like as the product: many brands, 100+ accounts each, posting daily and learning from results. The pipeline modules survive unchanged; everything around them is replaced by things built for volume.

Layer	Now	At scale
Front end	One local HTML page, no login.	Hosted web app with logins and teams; brands, drafts, approvals, calendar, results.
API and queue	app.py; a job list in memory, one job at a time.	Hosted API behind a durable job queue (Redis or SQS): thousands of jobs, retries, priorities.
Workers	One copy of make.py on a laptop.	Many copies of the same code as workers: profile and research, writing, and a render farm of ffmpeg workers.
Data	Files in out/brand/, clips in base/, a cache folder.	Postgres for brands, drafts, scores, posts, accounts, results. S3 for clips and videos. A searchable licensed clip library. Redis cache.
Model layer	Direct calls to Claude Opus 5 for every step.	A gateway with caching and cost tracking. Cheap models (Sonnet, Haiku or local) for extraction and filtering; Opus kept for drafts and critique. Prompts versioned and scored against real results.
Distribution	Not built. Download an MP4.	Scheduler decides which variant posts on which account at what slot. Posting service for the fleet. Analytics ingest: views, 3-second hold, completion, shares, comments.
Feedback loop	Six variants exist only to pick one post.	Variants become an experiment: post across accounts in the same window, shift slots toward hooks with the best hold rate after 48 hours (bandit allocation), re-weight the editor rubric by what performed, promote winning formats to templates.
Observability	Print statements in a terminal.	Logs, cost per video, failure alerts, model drift checks.

Build order: queue and workers first (same code, more of it), then Postgres and S3, then the hosted app, then posting and analytics, and the feedback loop last because it needs that posting data to exist.

Quality and reliability

How we know the output is good and the system is up, now and at scale

Content quality

Check	Now	At scale
On-brand	Editor scores whether the nouns and pains could only belong to this niche; drafts that could be swapped to another brand score low.	Same check, plus a per-brand voice profile learned from the posts that performed, and a human approval step before first post on a new brand.
Honesty	Every draft is checked against the brand profile; invented numbers or claims score 1 and cannot win. Research quotes must be exact substrings of the source.	Automated fact check against the profile and a claims list per brand; legal and platform policy rules encoded as hard fails; audit trail per post.
Taste	Style rules in a plain text file; ten brands reviewed by hand across niches.	A 20-plus brand eval set rated by humans. Any prompt or model change must beat the current version on it before it ships. Editor rubric re-weighted weekly by real 3-second hold rate.
Variants	Six per brand, best one chosen by score.	N per brand per day, allocated across accounts, winners promoted, losers retired. Performance, not opinion, decides.
Clip fit	Mood inferred from the text, ranked against a labelled library, reasons shown, human override.	Larger labelled library of creator footage; clip choice becomes one of the tested variables; low performers flagged automatically.

Reliability

Failure	Now	At scale
Site unreadable	Fallback to Claude web fetch; clear error if both fail.	Headless browser render as a third path; cached profiles so a flaky site does not block a daily run.
Model output malformed	One automatic JSON-repair retry.	Structured output schemas, retries with backoff, and a dead-letter queue for jobs that keep failing.
Search down	Falls back to evergreen only; the post still ships.	Multiple providers, cached research per niche, and a stale-but-usable policy.
No clip matches	Hard error, never a random clip; nothing lost, person picks one and rebuilds.	Larger library makes it rare; alert to a human editor; brand-level default clip set.
Render fails	Writing is already saved; render is retried alone.	Render farm with retries; corrupt files detected by decoding every output before it is marked done.
Cost or rate limits	Under 1 pound per brand; no limits hit.	Per-brand budgets, cheap models for cheap steps, batch API overnight, alerts on cost per video.
Laptop off	Demo goes down.	Hosted, multi-worker, health checks, automatic job recovery.

Rule for the whole loop: if it can be done with counters, do that; if it needs a model, start with a linear one; use a large language model only to explain what the numbers already showed.